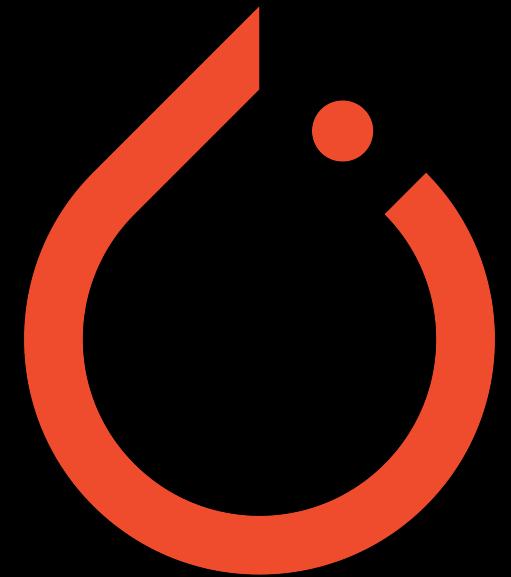




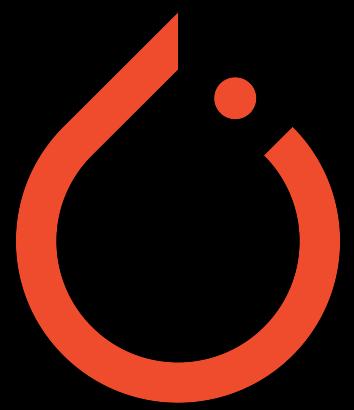
# “What is PyTorch?”



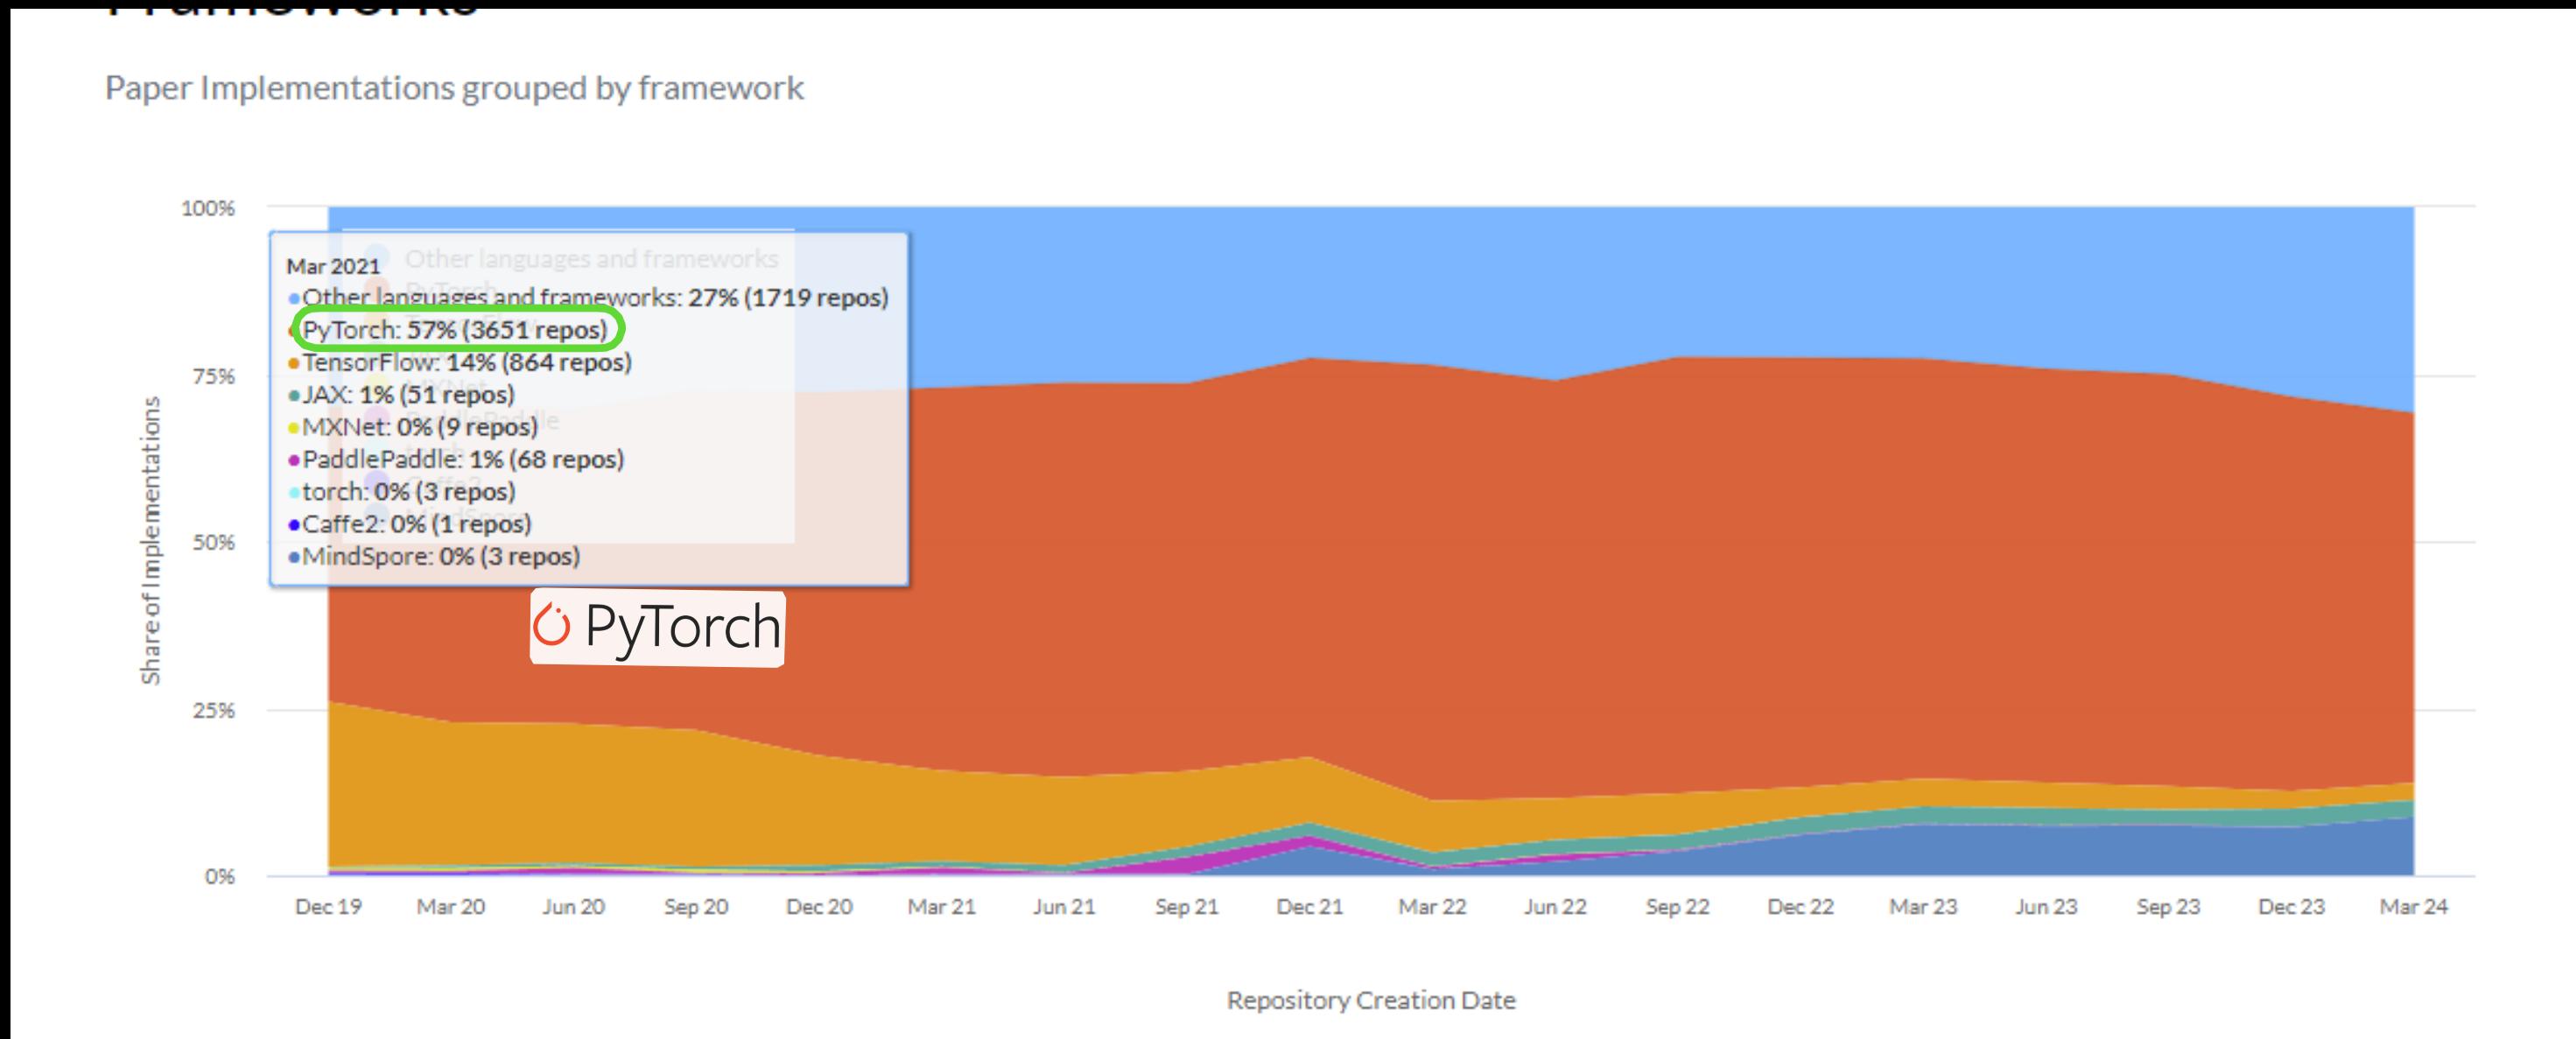
# What is PyTorch?



- Most popular research deep learning framework\*



# Why PyTorch?



Research favorite

# What is PyTorch?

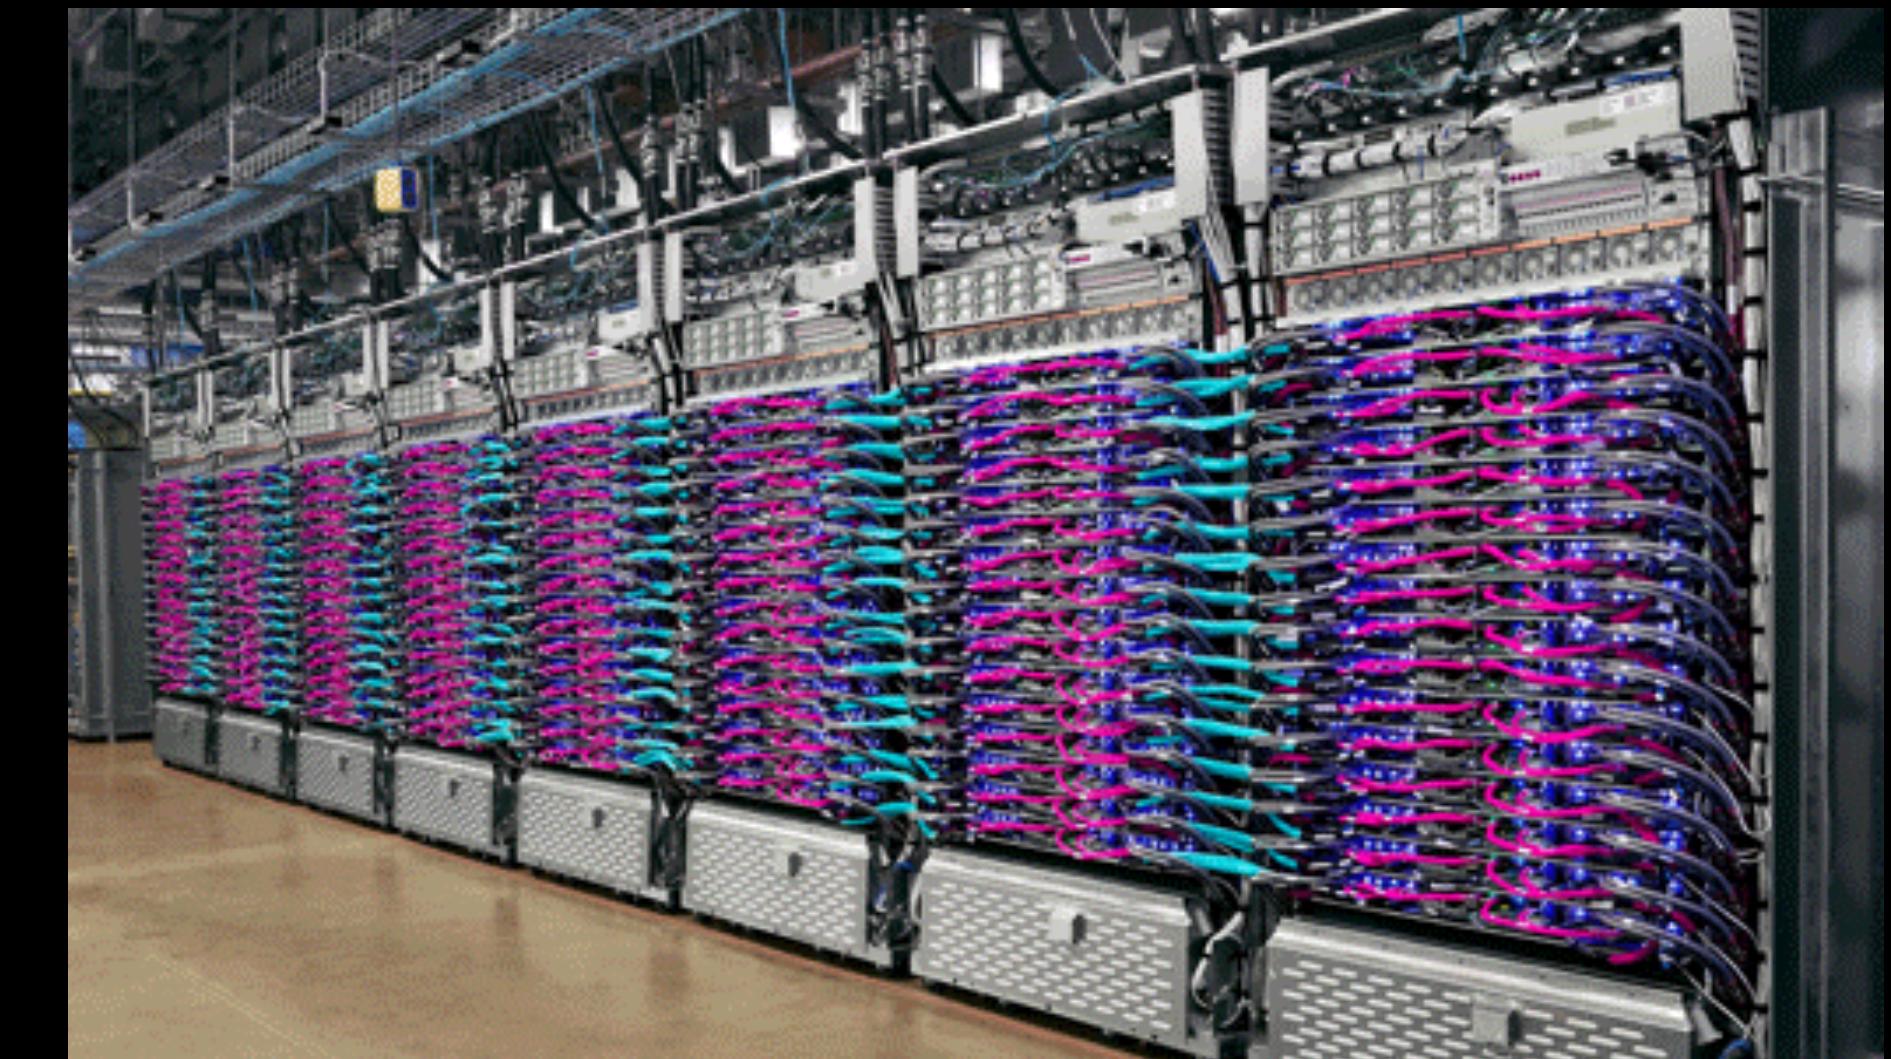


- Most popular research deep learning framework\*
- Write fast deep learning code in Python (able to run on a GPU/many GPUs)

# What is a GPU/TPU?



**GPU (Graphics Processing Unit)**



**TPU (Tensor Processing Unit)**

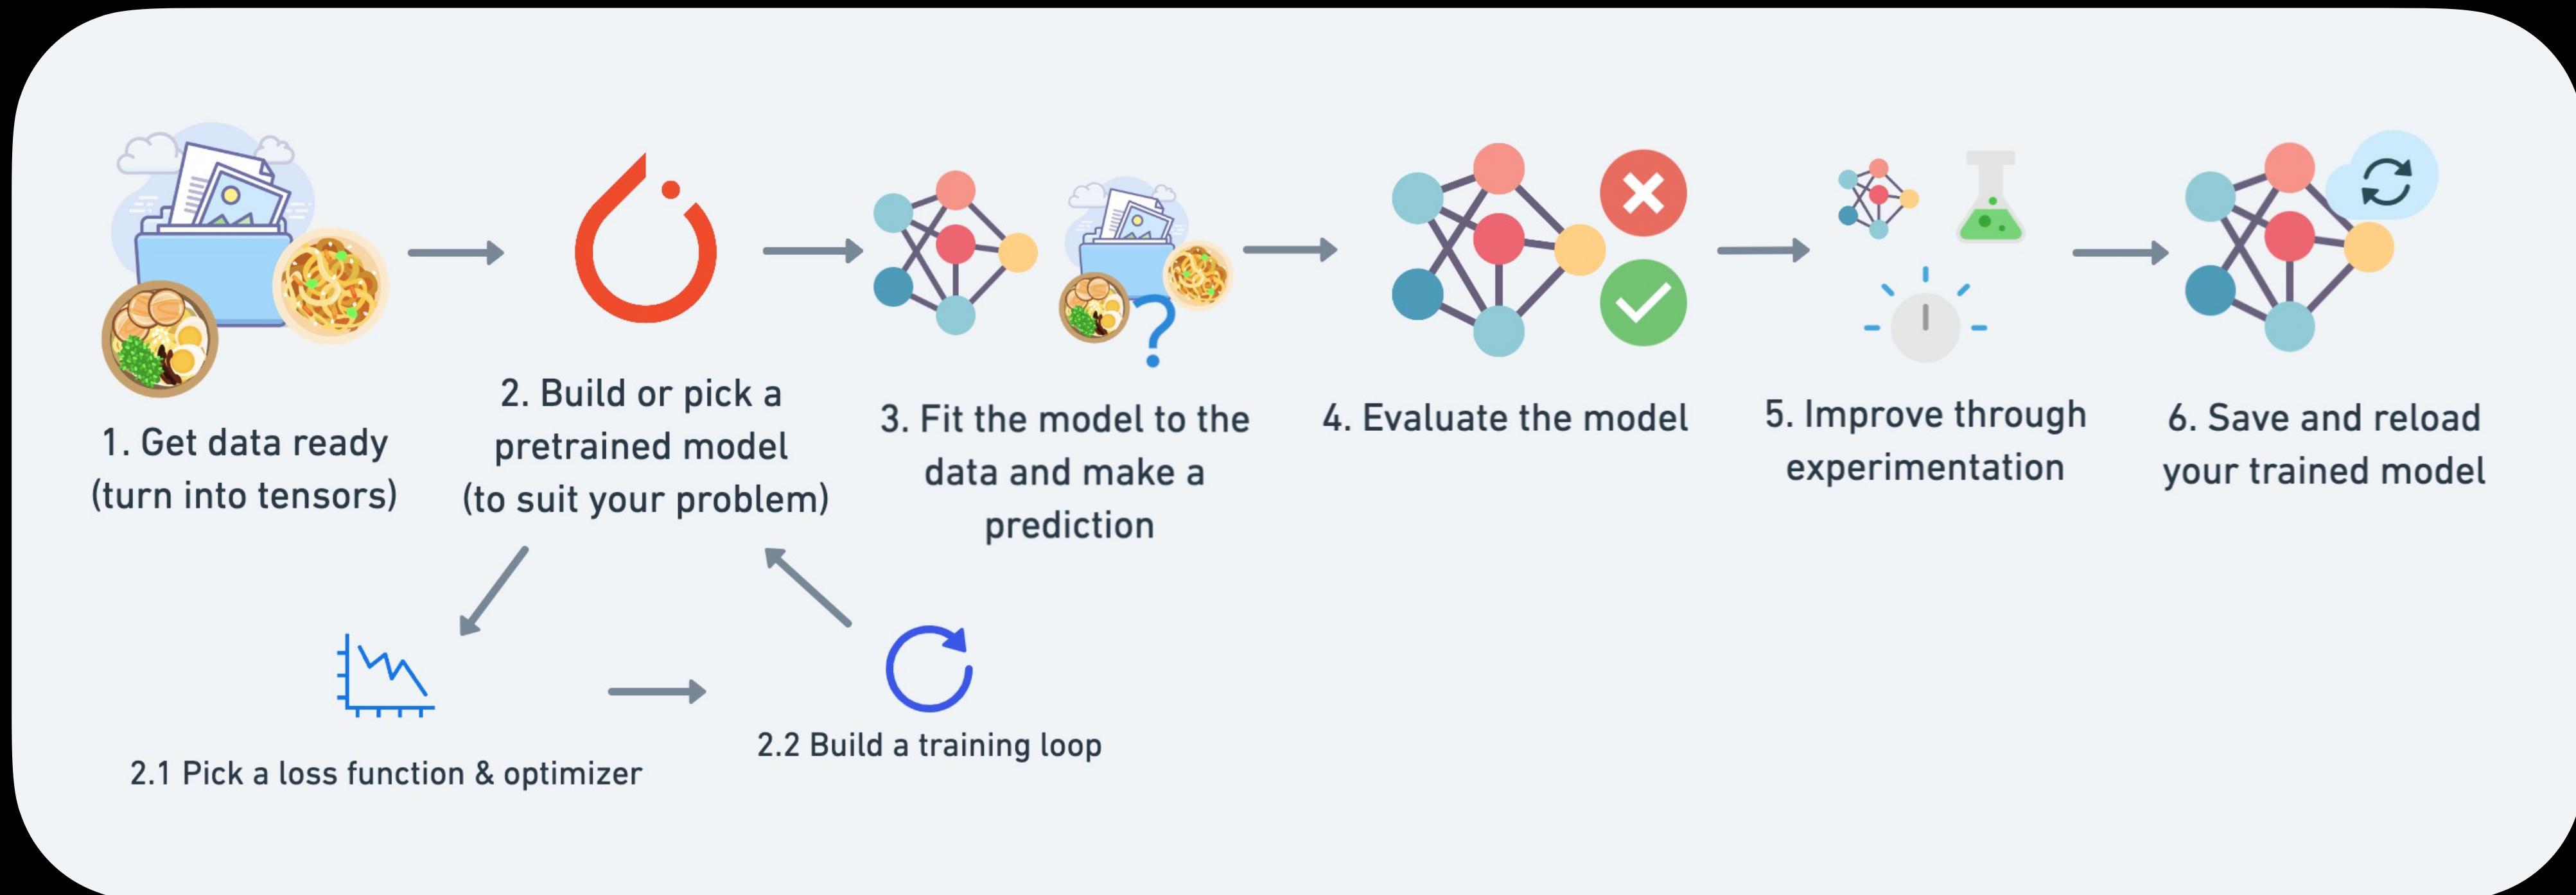
# What is PyTorch?



- Most popular research deep learning framework\*
- Write fast deep learning code in Python (able to run on a GPU/many GPUs)
- Able to access many pre-built deep learning models (Torch Hub/torchvision.models, Hugging Face)
- Whole stack: preprocess data, model data, deploy model in your application/cloud
- Originally designed and used in-house by Facebook/Meta (now open-source and used by companies such as Tesla, Microsoft, OpenAI)

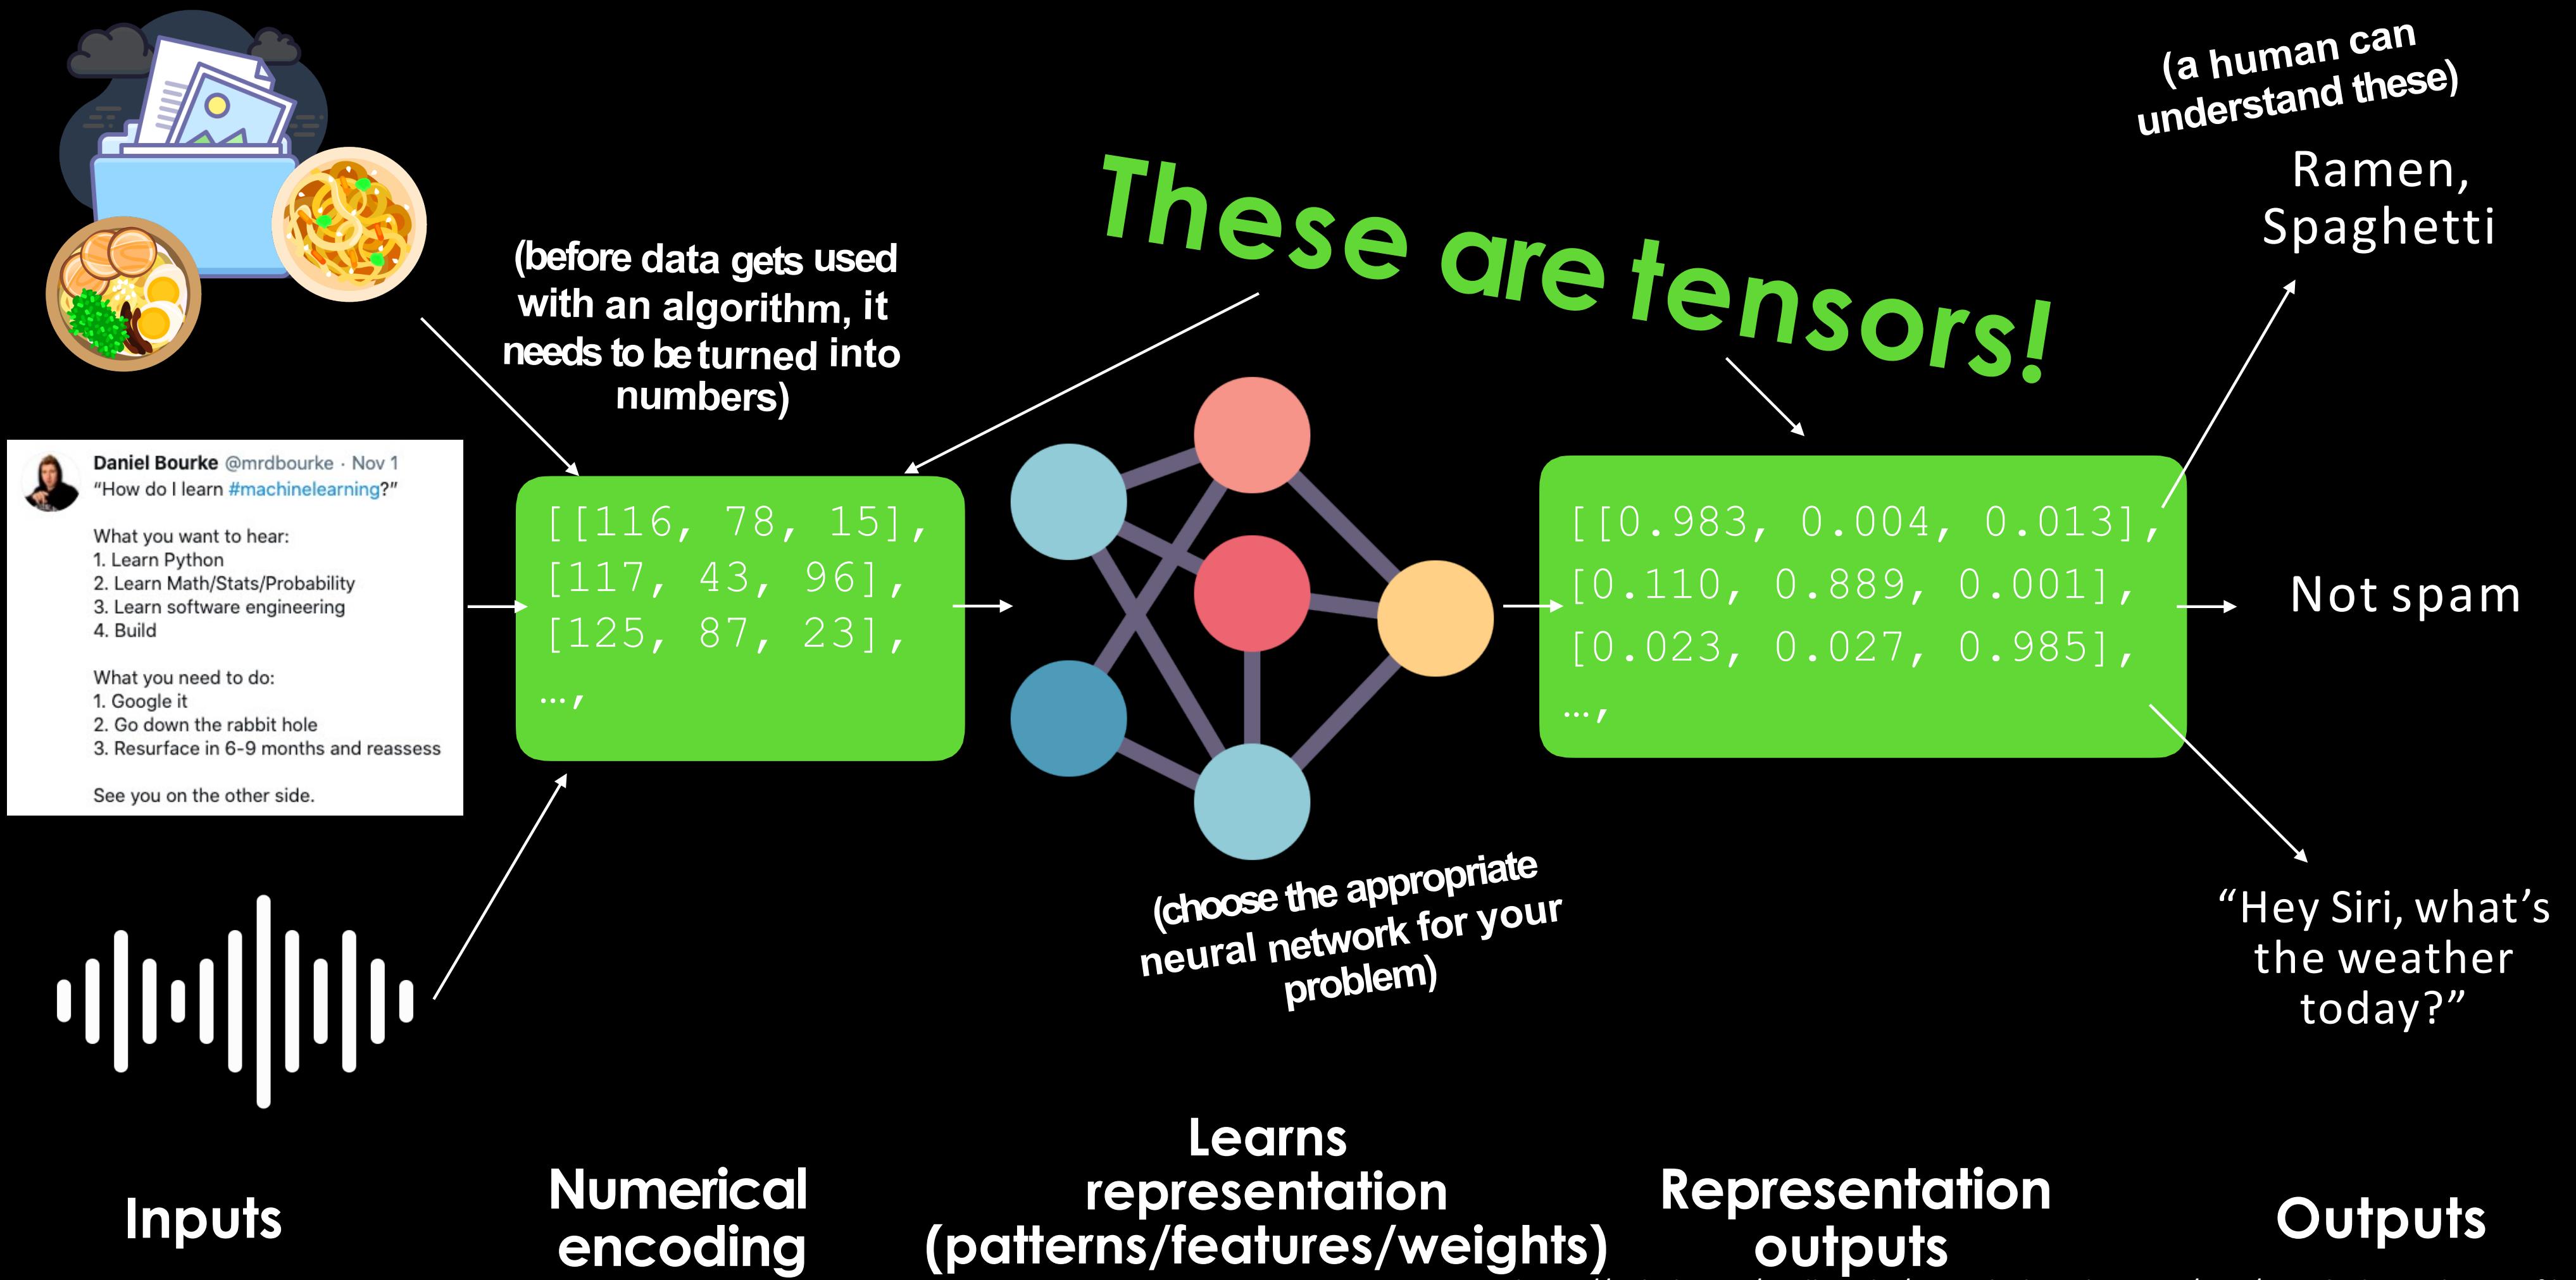
# What we're going to cover

## A PyTorch workflow (one of many)



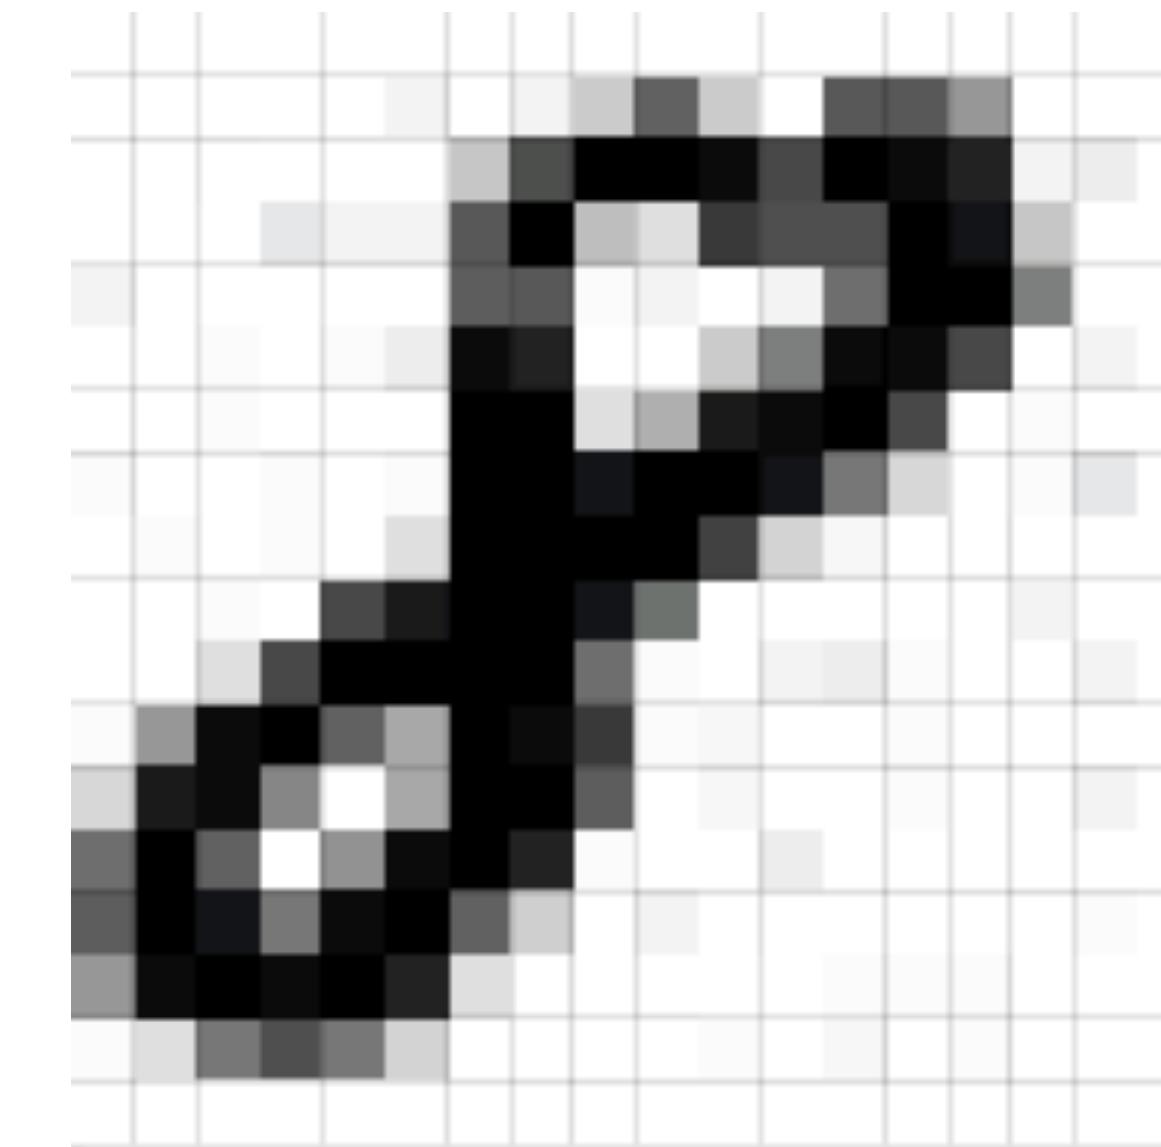
# “What is a tensor?”

# Neural Networks

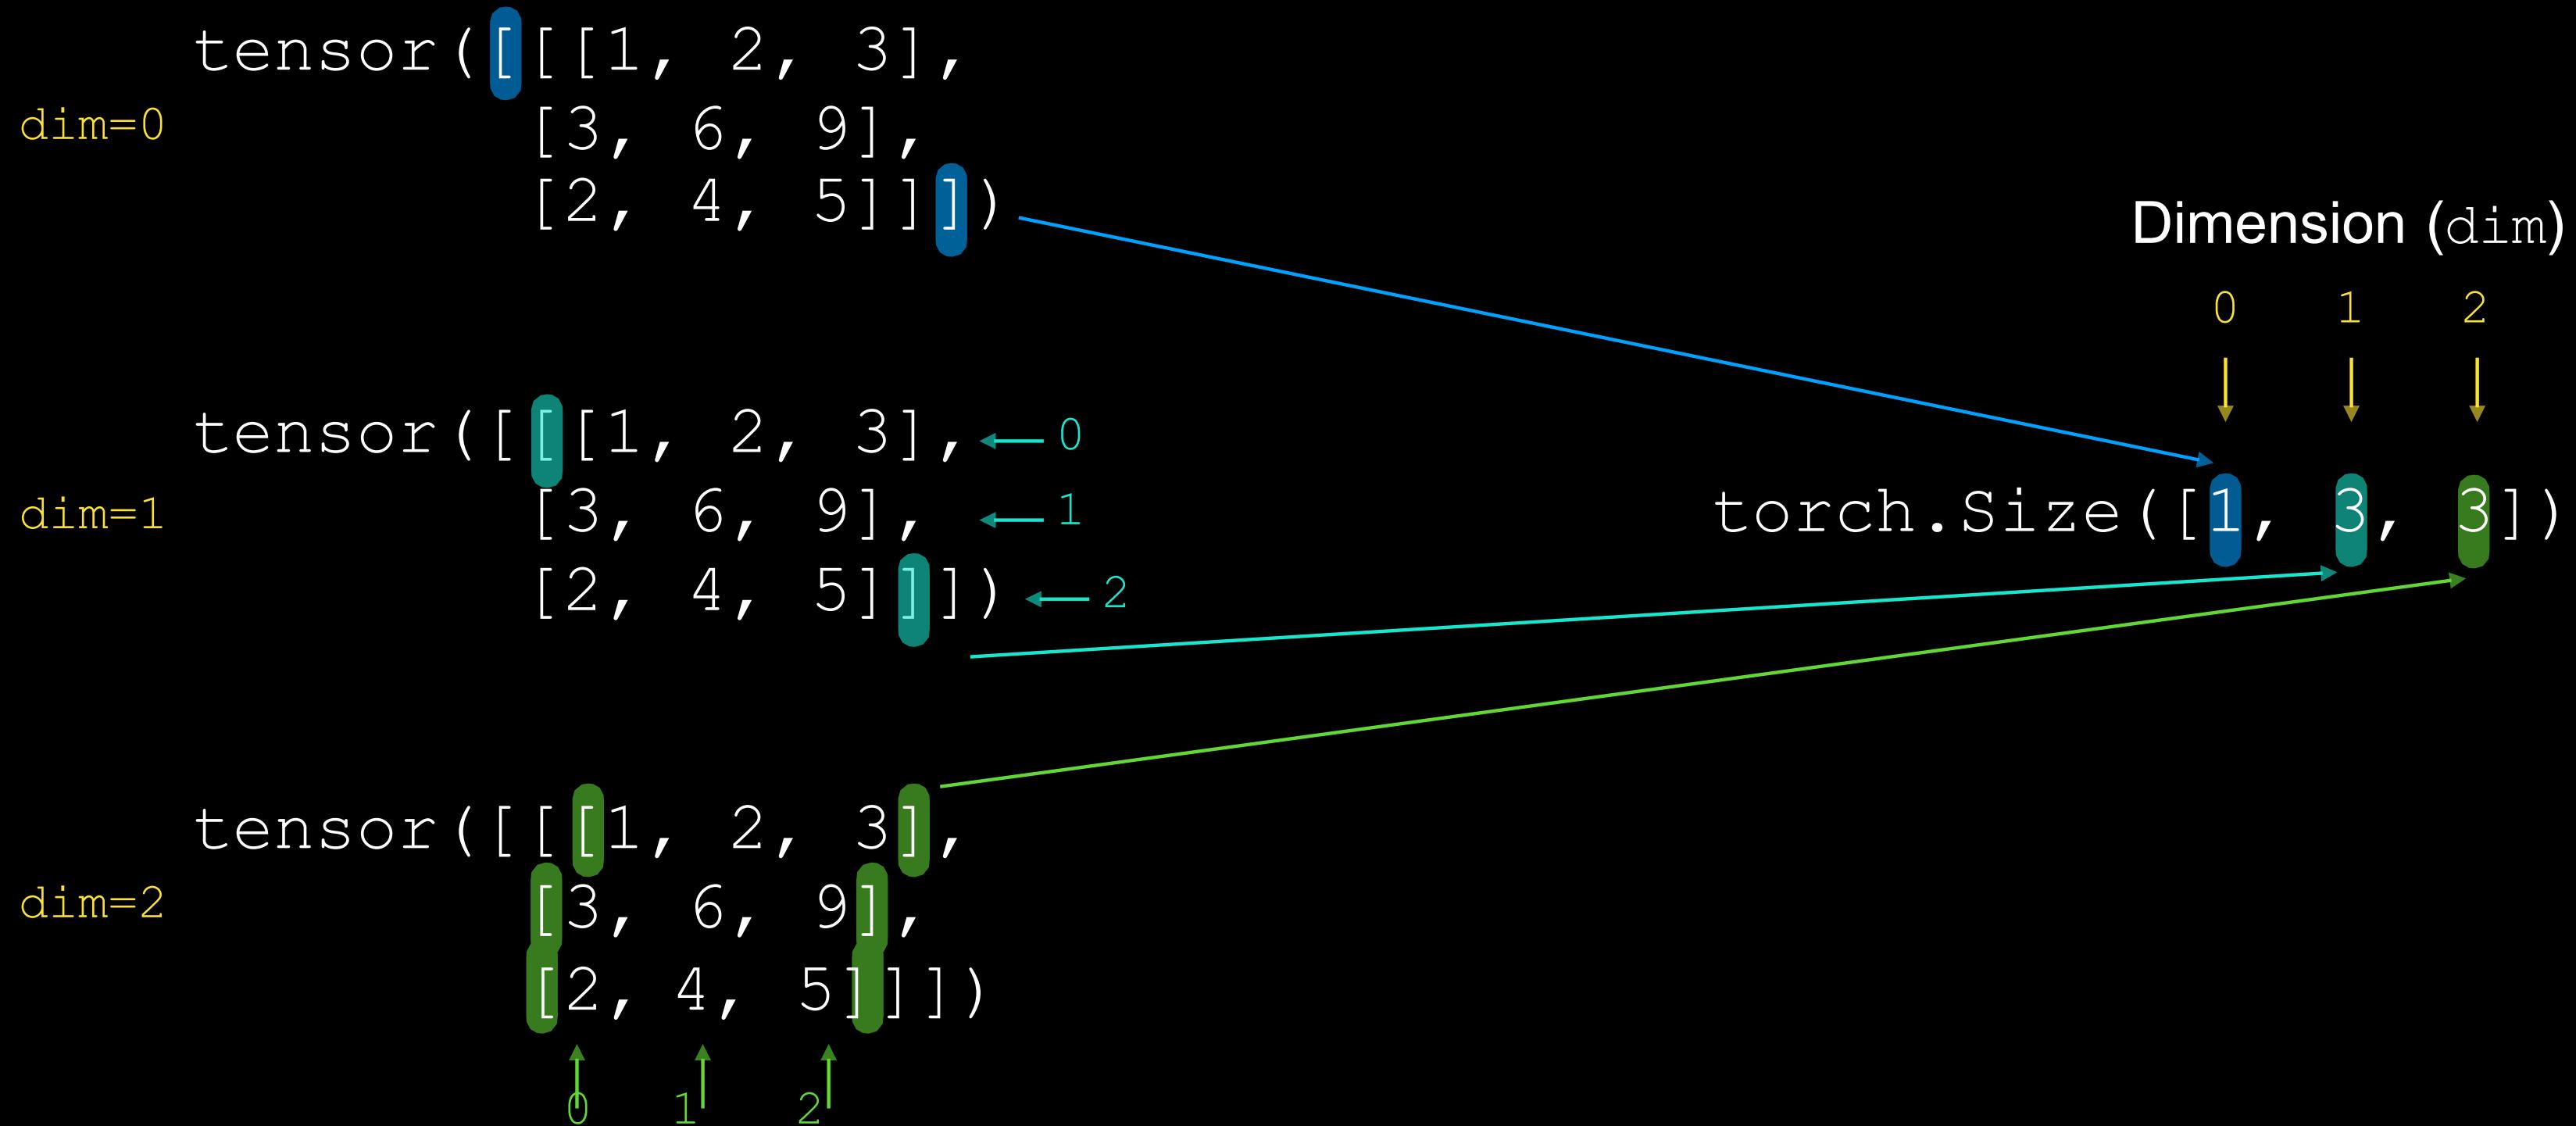


**These are tensors!**

An Image is  
a matrix of  
pixel values

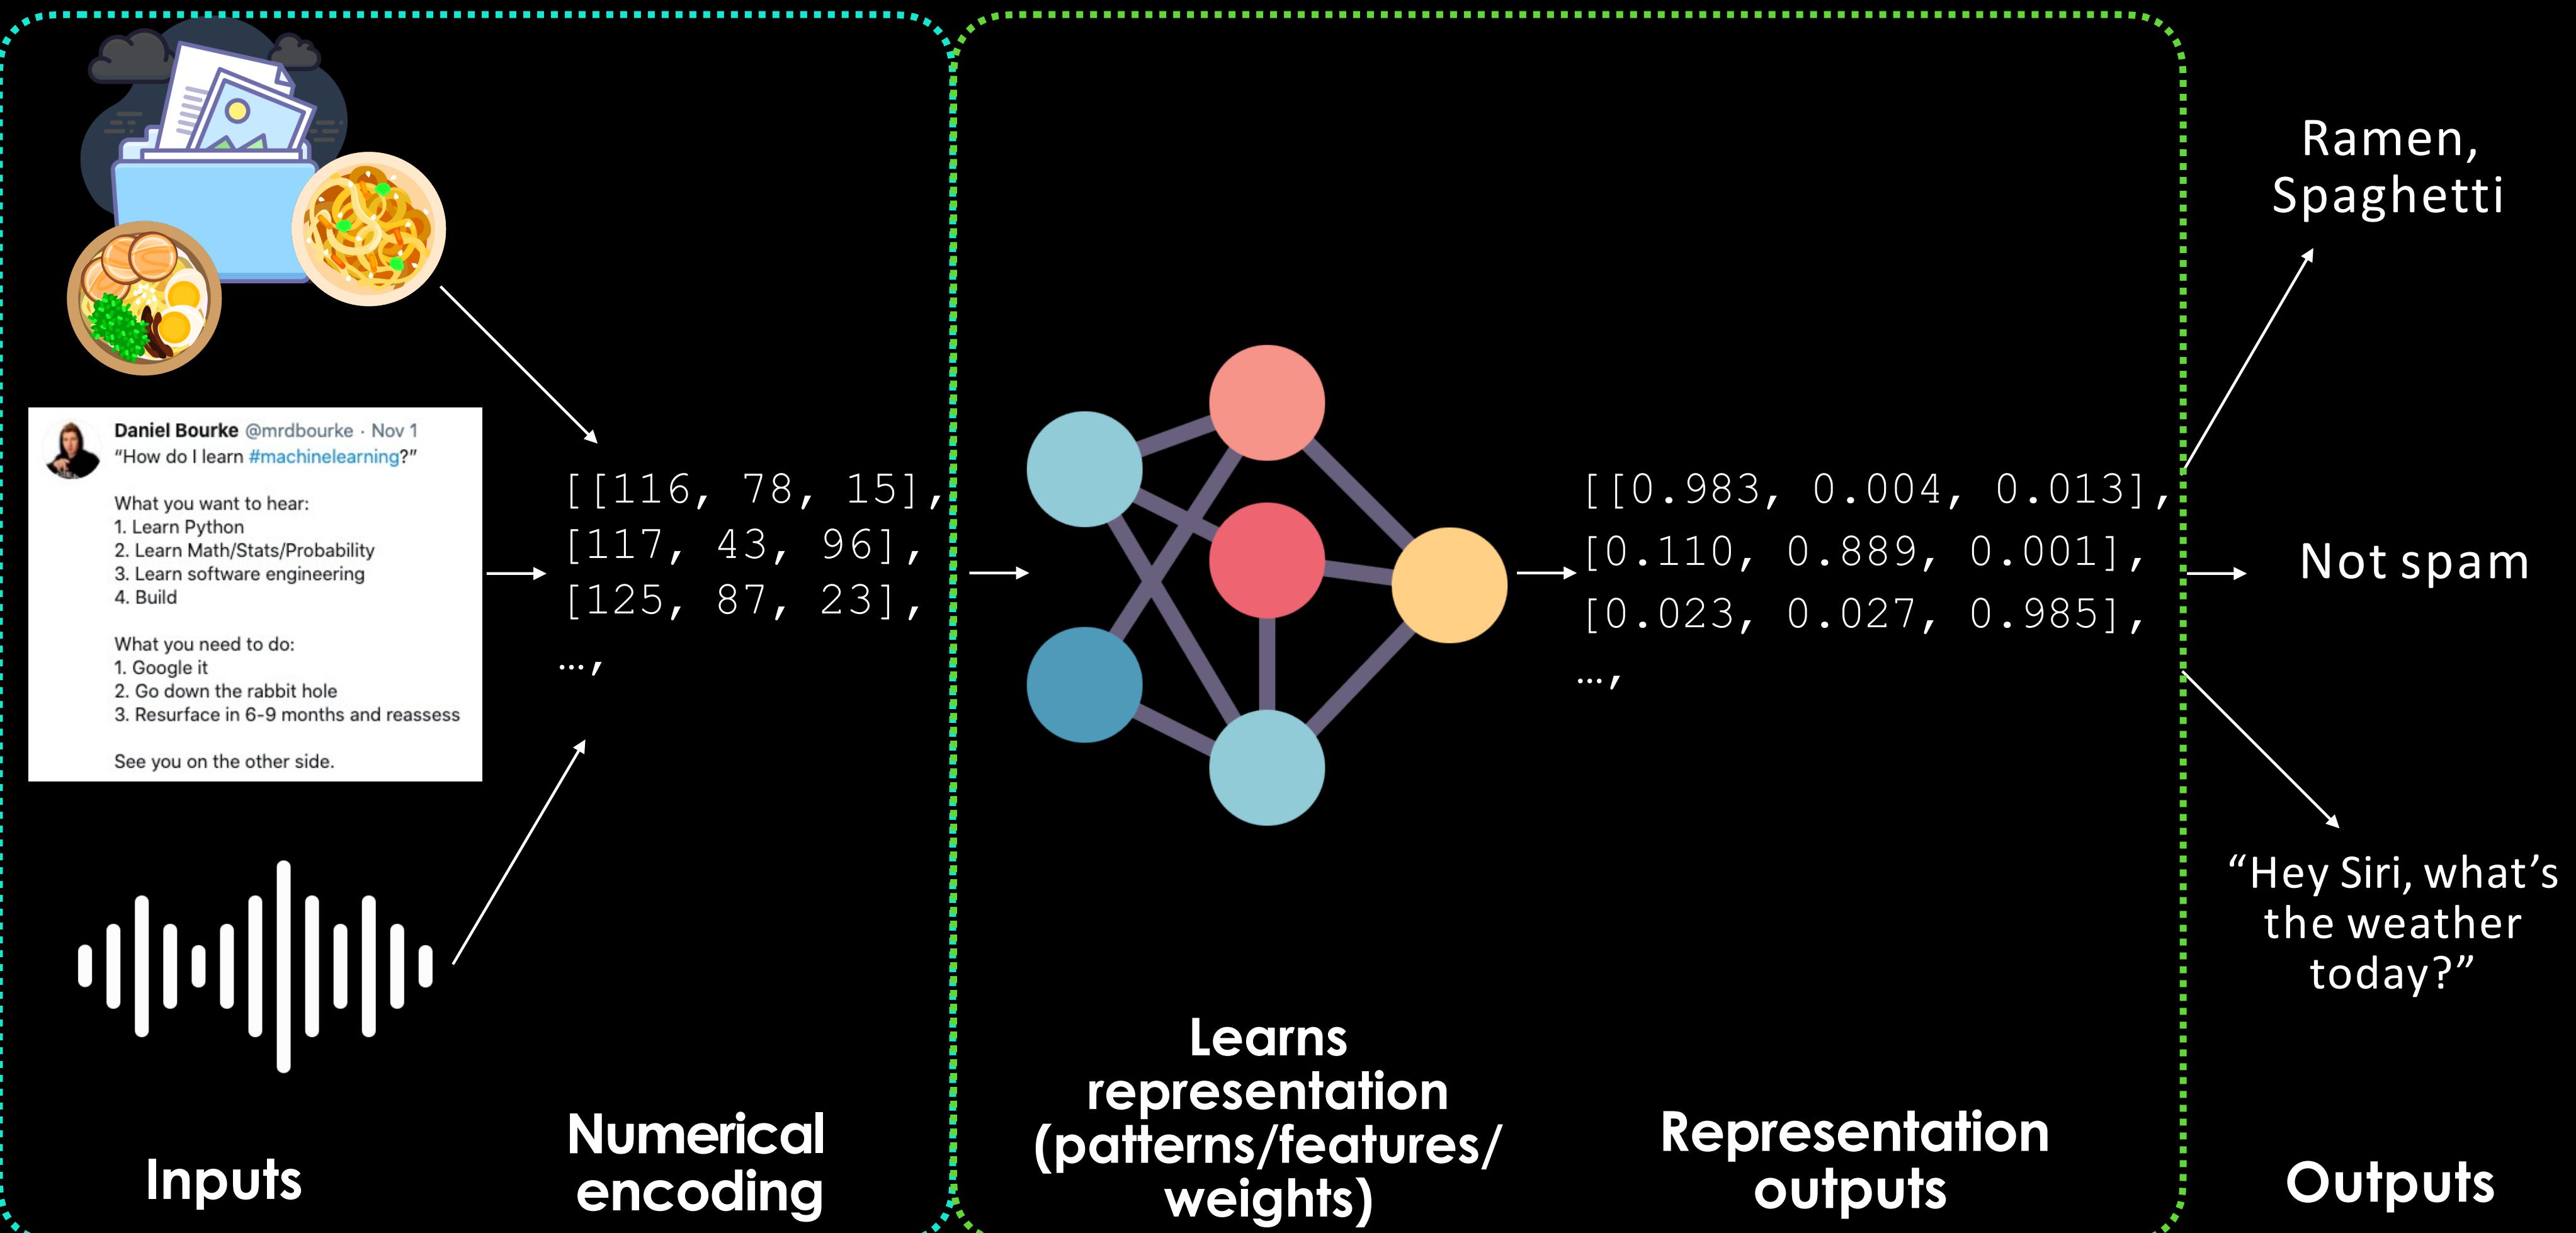


# Tensor dimensions

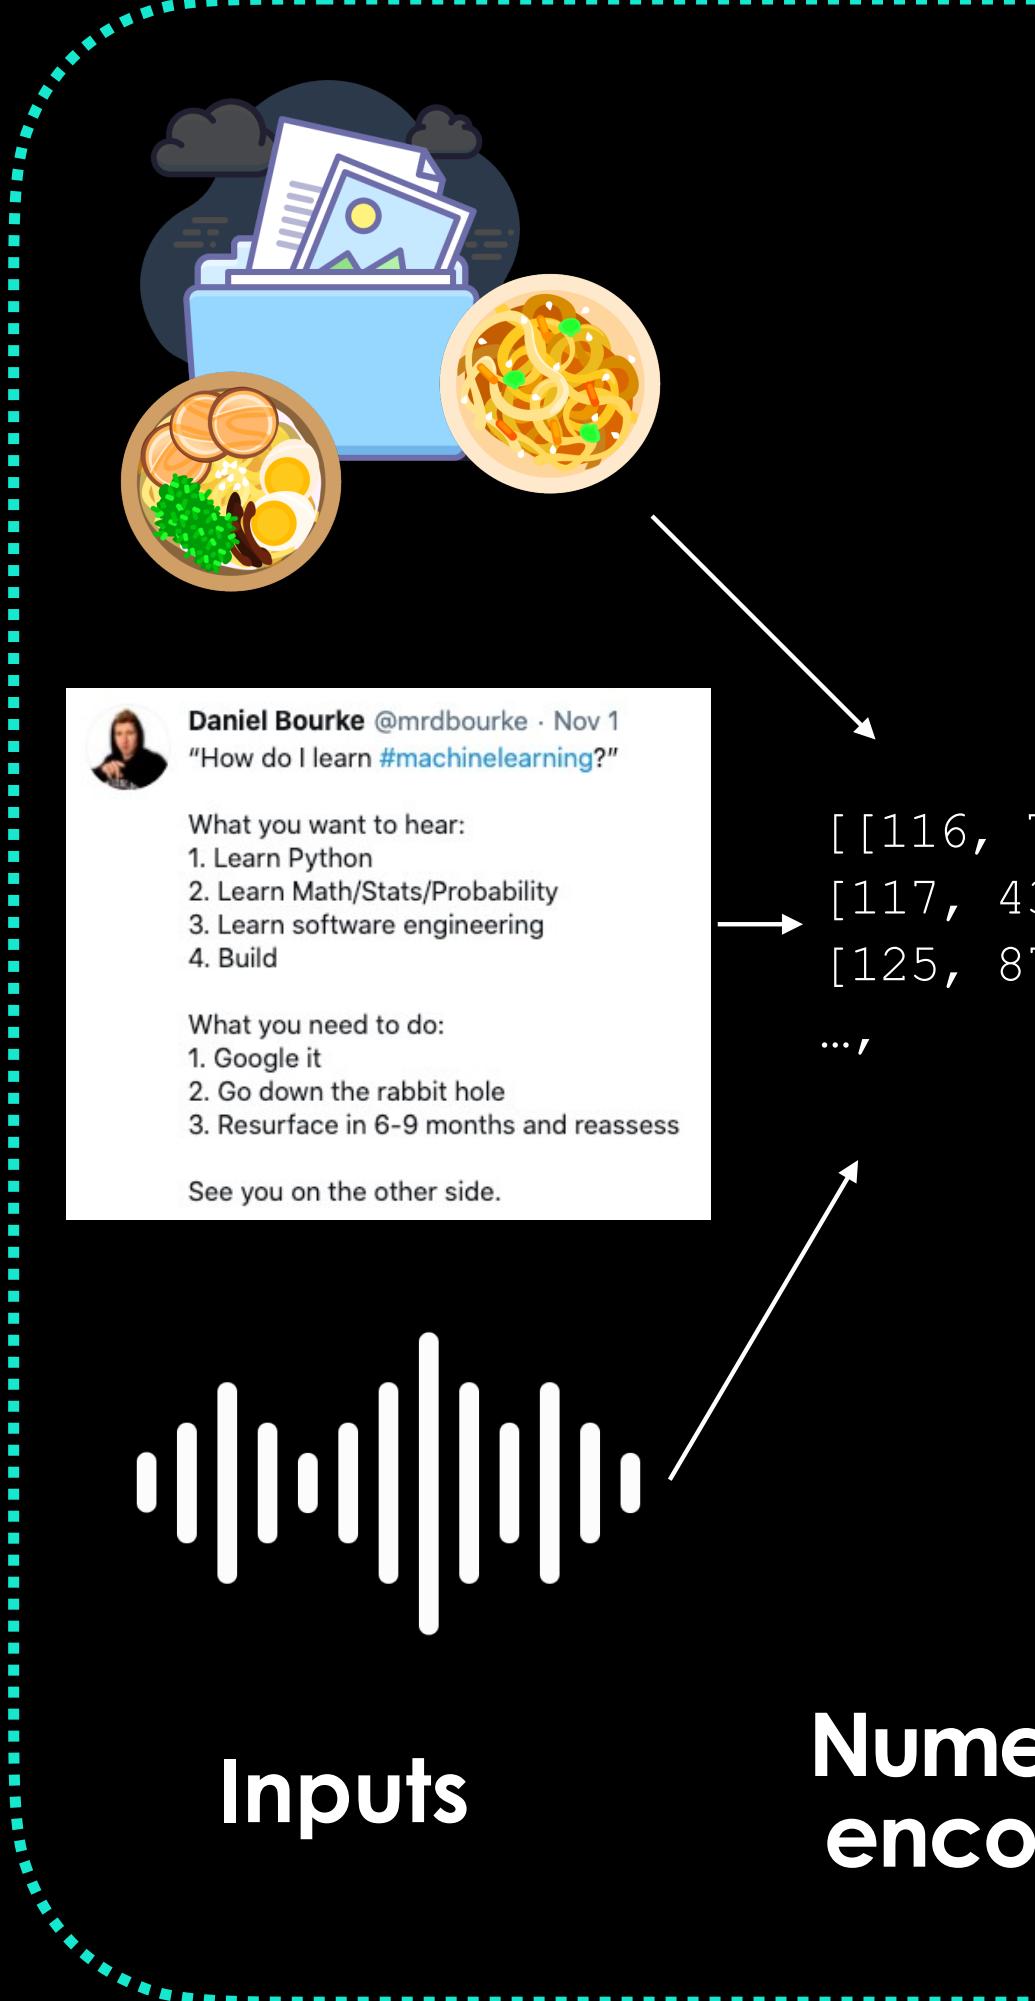


# 01 tensors.ipynb

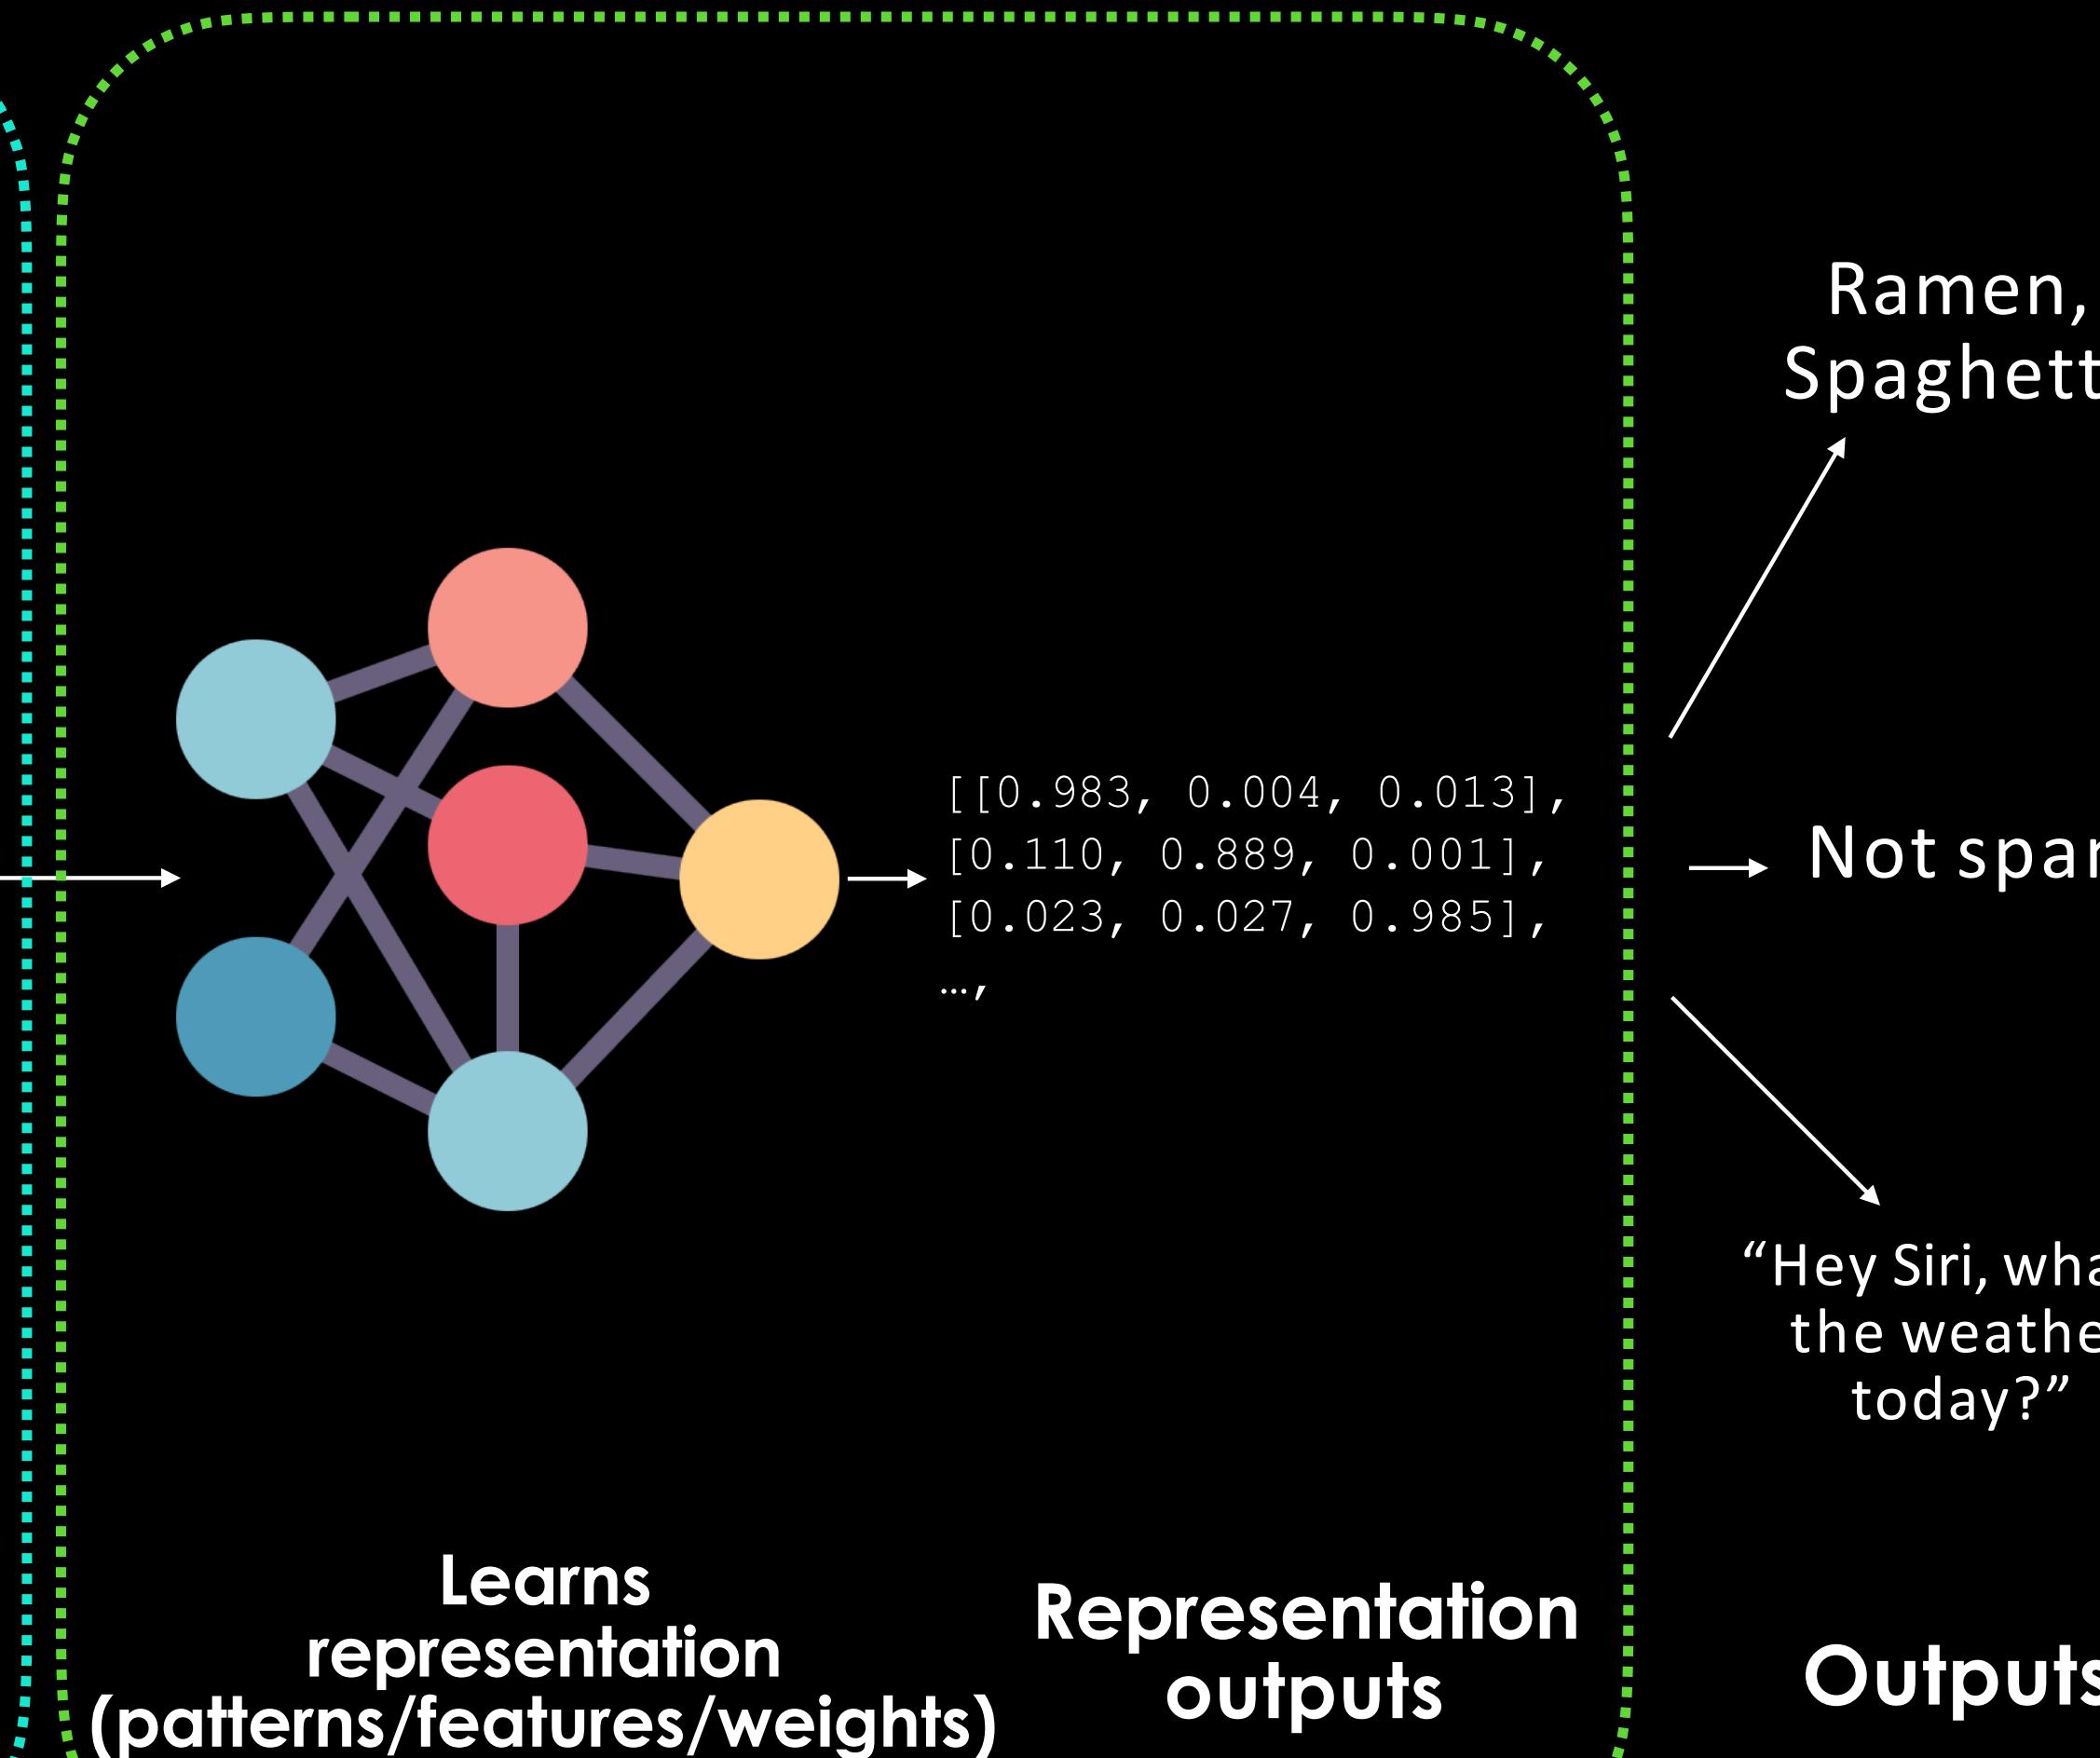
# Machine learning: a game of two parts



## Part 1: Turn data into numbers



## Part 2: Build model to learn patterns in numbers



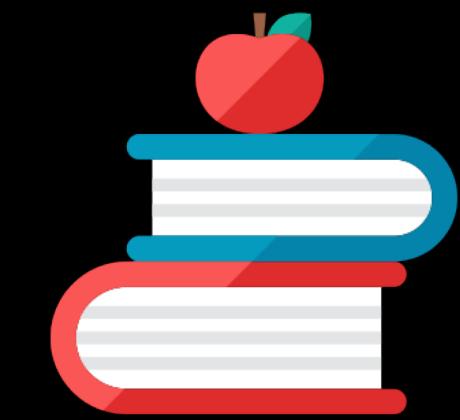
# Three datasets

(possibly the most important concept in machine learning...)

Model learns patterns from here



Course materials  
(**training set**)



Practice exam  
(**validation set**)  
Tune model patterns



Final exam  
(**test set**)  
See if the model is ready for the wild

## Generalization

The ability for a machine learning model to perform well on data it hasn't seen before.

02 dataset and dataloader.ipynb

03 transforms.ipynb

# Supervised learning (overview)

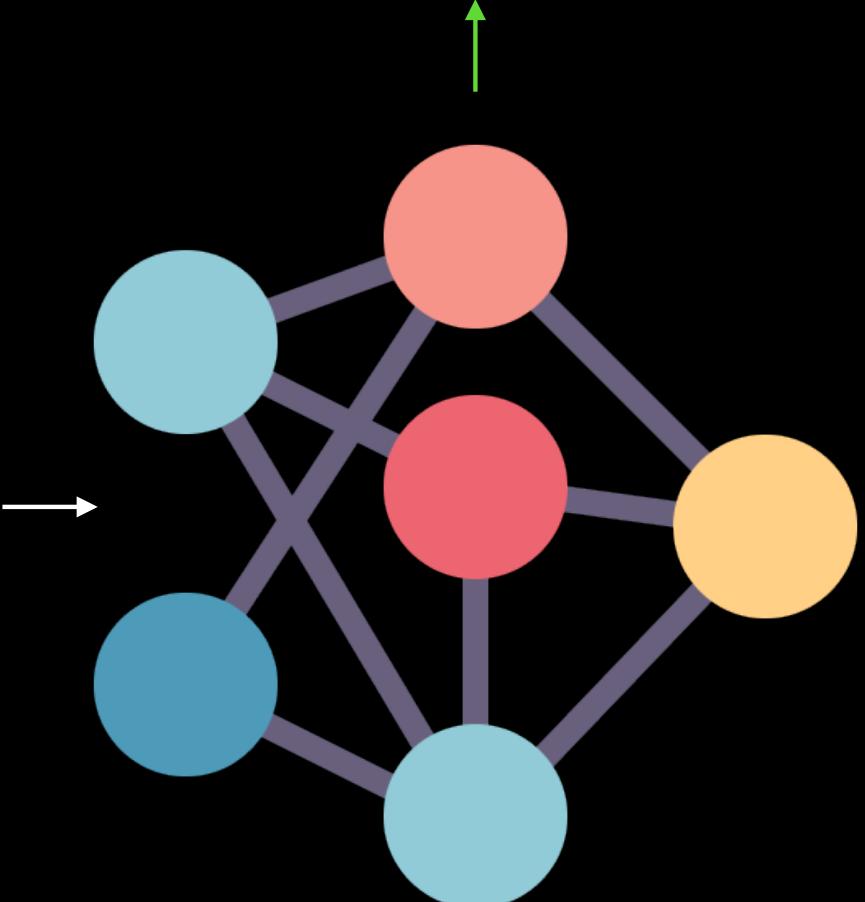


**4. Repeat with more examples**

**1. Initialise with random weights (only at beginning)**

$\begin{bmatrix} [0.092, 0.210, 0.415], \\ [0.778, 0.929, 0.030], \\ [0.019, 0.182, 0.555], \end{bmatrix}$

...



$\rightarrow \begin{bmatrix} [116, 78, 15], \\ [117, 43, 96], \\ [125, 87, 23], \end{bmatrix}$

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

```
1 # Create a linear regression model in PyTorch
2 class LinearRegressionModel(nn.Module):
3     def __init__(self):
4         super().__init__()
5
6         # Initialize model parameters
7         self.weights = nn.Parameter(torch.randn(1,
8             requires_grad=True,
9             dtype=torch.float
10        ))
11
12        self.bias = nn.Parameter(torch.randn(1,
13            requires_grad=True,
14            dtype=torch.float
15        ))
16
17        # forward() defines the computation in the model
18    def forward(self, x: torch.Tensor) -> torch.Tensor:
19        return self.weights * x + self.bias
```

Subclass `nn.Module`  
(this contains all the building blocks for neural networks)

Initialise **model parameters** to be used in various computations (these could be different layers from `torch.nn`, single parameters, hard-coded values or functions)

`requires_grad=True` means PyTorch will track the gradients of this specific parameter for use with `torch.autograd` and gradient descent (for many `torch.nn` modules, `requires_grad=True` is set by default)

Any subclass of `nn.Module` needs to override `forward()`  
(this defines the forward computation of the model)

```
1 # Create a linear regression model in PyTorch
2 class LinearRegressionModel(nn.Module):
3     def __init__(self):
4         super().__init__()
5
6         # Initialize model parameters
7         self.weights = nn.Parameter(torch.randn(1,
8             requires_grad=True,
9             dtype=torch.float
10            ))
11
12         self.bias = nn.Parameter(torch.randn(1,
13             requires_grad=True,
14             dtype=torch.float
15            ))
16
17     # forward() defines the computation in the model
18     def forward(self, x: torch.Tensor) -> torch.Tensor:
19         return self.weights * x + self.bias
```

## Linear regression model with nn.Parameter

```
1 # Create a linear regression model in PyTorch with nn.Linear
2 class LinearRegressionModel(nn.Module):
3     def __init__(self):
4         super().__init__()
5         # Use nn.Linear() for creating the model parameters
6         self.linear_layer = nn.Linear(in_features=1,
7                                         out_features=1)
8
9     # forward() defines the computation in the model
10    def forward(self, x: torch.Tensor) -> torch.Tensor:
11        return self.linear_layer(x)
```

## Linear regression model with nn.Linear

# PyTorch essential neural network building modules

PyTorch module	What does it do?
<u><a href="#">torch.nn</a></u>	Contains all of the building blocks for computational graphs (essentially a series of computations executed in a particular way).
<u><a href="#">torch.nn.Module</a></u>	The base class for all neural network modules, all the building blocks for neural networks are subclasses. If you're building a neural network in PyTorch, your models should subclass <code>nn.Module</code> . Requires a <code>forward()</code> method be implemented.
<u><a href="#">torch.optim</a></u>	Contains various optimization algorithms (these tell the model parameters stored in <code>nn.Parameter</code> how to best change to improve gradient descent and in turn reduce the loss).
<u><a href="#">torch.utils.data.Dataset</a></u>	Represents a map between key (label) and sample (features) pairs of your data. Such as images and their associated labels.
<u><a href="#">torch.utils.data.DataLoader</a></u>	Creates a Python iterable over a torch Dataset (allows you to iterate over your data).

torchvision.transforms  
torch.utils.data.Dataset

torch.utils.data.DataLoader



1. Get data ready  
(turn into tensors)



2. Build or pick a  
pretrained model  
(to suit your problem)



3. Fit the model to the  
data and make a  
prediction

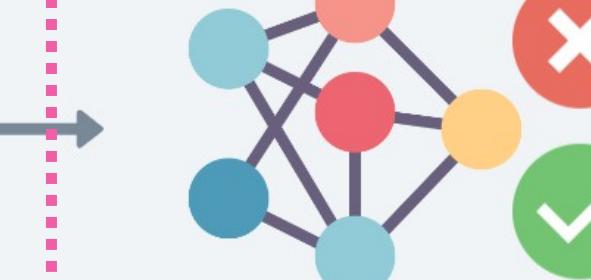


2.1 Pick a **loss function** & **optimizer**



2.2 Build a training loop

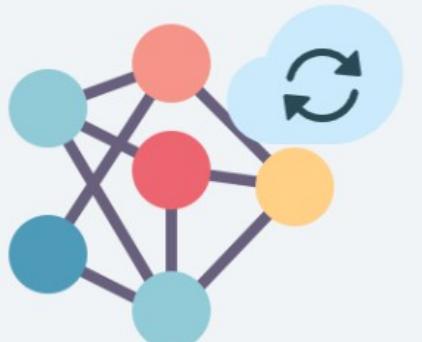
torchmetrics



4. Evaluate the model



5. Improve through  
experimentation



6. Save and reload  
your trained model

torch.optim

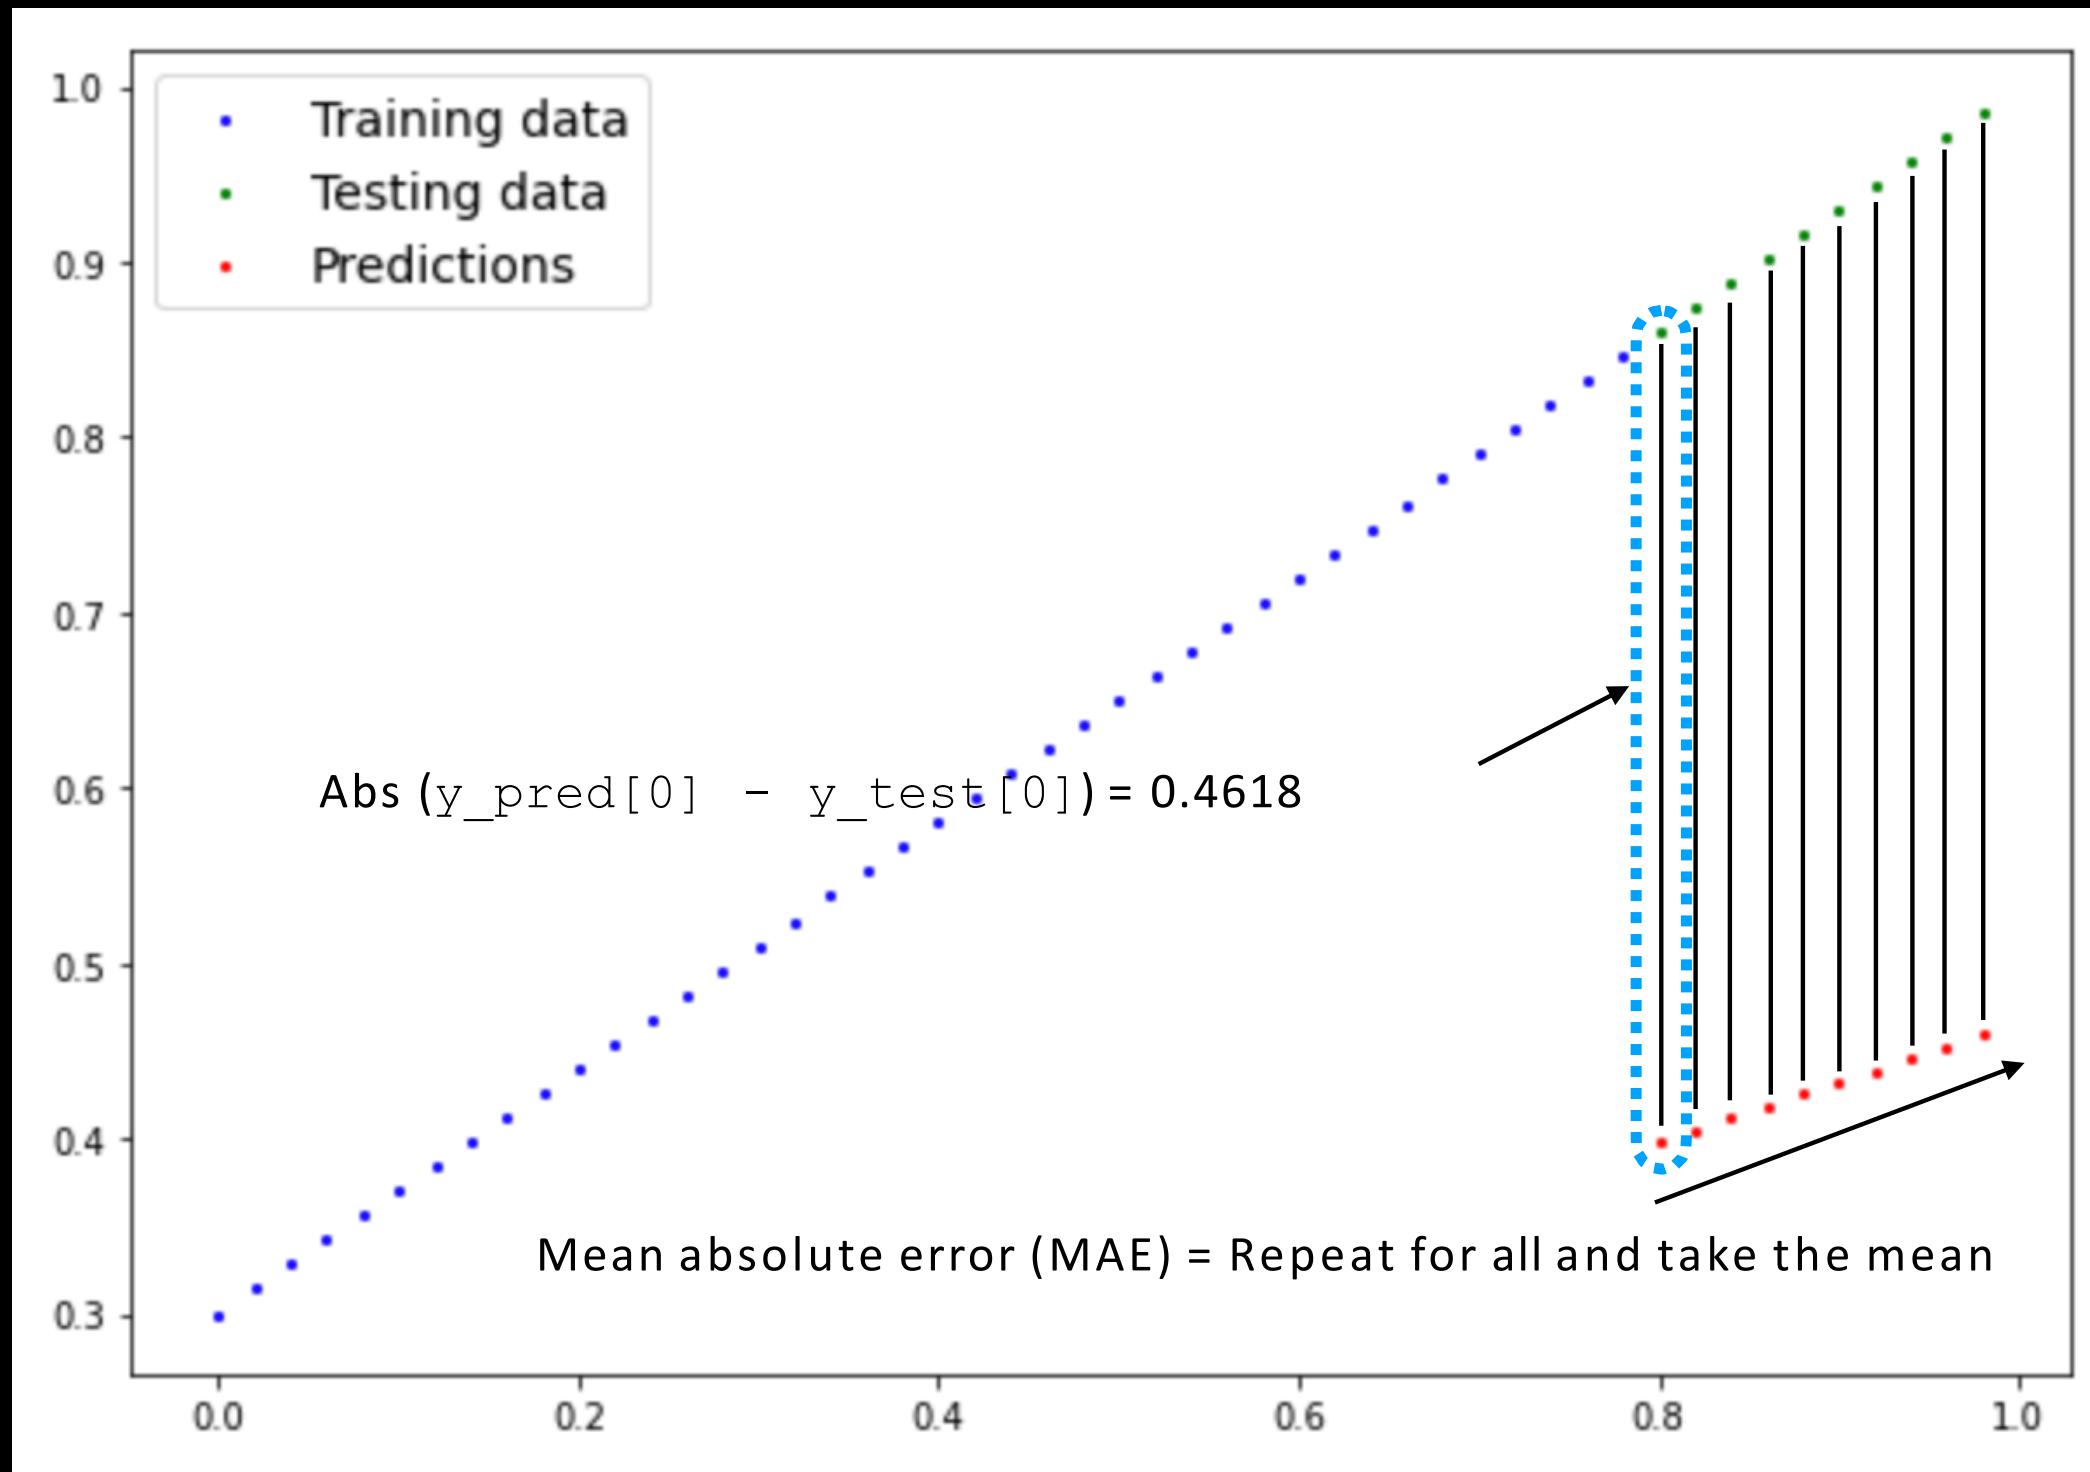
torch.nn

torch.nn.Module

torchvision.models

torch.utils.tensorboard

# Mean absolute error (MAE)



```
MAE_loss = torch.mean(torch.abs(y_pred-y_test))  
or  
MAE_loss = torch.nn.L1Loss
```

# Binary Cross Entropy (torch.nn.BCELoss)

# Cross Entropy Loss (torch.nn.CrossEntropyLoss)

04 building the neural network.ipynb  
05 autograd.ipynb

# PyTorch training loop

```
1 # Pass the data through the model for a number of epochs (e.g. 100)
2 for epoch in range(epochs):
3
4     # Put model in training mode (this is the default state of a model)
5     model.train()
6
7     # 1. Forward pass on train data using the forward() method inside
8     y_pred = model(X_train)
9
10    # 2. Calculate the loss (how different are the model's predictions to the true values)
11    loss = loss_fn(y_pred, y_true)
12
13    # 3. Zero the gradients of the optimizer (they accumulate by default)
14    optimizer.zero_grad()
15
16    # 4. Perform backpropagation on the loss
17    loss.backward()
18
19    # 5. Progress/step the optimizer (gradient descent)
20    optimizer.step()
```

Note: all of this can be turned into a function

Pass the data through the model for a number of **epochs** (e.g. 100 for 100 passes of the data)

Pass the data through the model, this will perform the `forward()` method located within the `model` object

Calculate the **loss value** (how wrong the model's predictions are)

Zero the **optimizer gradients** (they accumulate every epoch, zero them to start fresh each forward pass)

Perform **backpropagation** on the loss function (compute the gradient of every parameter with `requires_grad=True`)

Step the **optimizer** to update the model's parameters with respect to the gradients calculated by `loss.backward()`

# PyTorch testing loop

```
1 # Setup empty lists to keep track of model progress
2 epoch_count = []
3 train_loss_values = []
4 test_loss_values = []
5
6 # Pass the data through the model for a number of epochs (e.g. 100) pochs):
7 for epoch in range(epochs):
8
9     ### Training loop code here #####
10
11    ### Testing starts #####
12
13    # Put the model in evaluation mode
14    model.eval()
15
16    # Turn on inference mode context manager
17    with torch.inference_mode():
18        # 1. Forward pass on test data
19        test_pred = model(X_test)
20
21        # 2. Calculate loss on test data
22        test_loss = loss_fn(test_pred, y_test)
23
24    # Print out what's happening every 10 epochs
25    if epoch % 10 == 0:
26        epoch_count.append(epoch)
27        train_loss_values.append(loss)
28        test_loss_values.append(test_loss)
29        print(f"Epoch: {epoch} | MAE Train Loss: {loss} | MAE Test Loss: {test_loss} ")
```

Note: all of this can be turned into a function

Create empty lists for storing useful values (helpful for tracking model progress)

Tell the model we want to evaluate rather than train (this turns off functionality used for training but not evaluation)

Turn on `torch.inference_mode()` context manager to disable functionality such as gradient tracking for inference (gradient tracking not needed for inference) (**faster performance!**)

Pass the test data through the model (this will call the model's implemented `forward()` method)

Calculate the test loss value (how wrong the model's predictions are on the test dataset, lower is better)

Display information outputs for how the model is doing during training/testing every ~10 epochs (note: what gets printed out here can be adjusted for specific problems)

```
1 # Setup optimization loop(s)
2 epochs = 10000
3
4 ### Train time!
5 # Loop through the epochs
6 for epoch in range(epochs):
7     # Set the model to train mode (this is the default)
8     model.train()
9
10    # 1. Do the forward pass
11    y_pred = model(x_train)
12
13    # 2. Calculate the loss (how wrong the model is)
14    loss = loss_fn(y_pred, y_train)
15
16    # 3. Zero the optimizer gradients (they accumulate by default)
17    optimizer.zero_grad()
18
19    # 4. Perform backpropagation (with respect to the model's parameters)
20    loss.backward()
21
22    # 5. Step the optimizer (gradient descent)
23    optimizer.step()
24
25 ### Test time!
26 # Set the model to eval mode (this turns off settings not needed for testing)
27 model.eval()
28 # Turn on inference mode context manager (removes even more things not needed for inference)
29 with torch.inference_mode():
30     # 1. Do the forward pass
31     test_pred = model(x_test)
32
33     # 2. Calculate the loss
34     test_loss = loss_fn(test_pred, y_test)
35
36 # Print out what's happenin'!
37 print(f"Epoch: {epoch} | Train loss: {loss:.4f} | Test loss: {test_loss:.4f}")
```

```
1 # Train function
2 def train_step(model, loss_fn, optimizer, data, labels):
3     # Turn on train mode (this is default but we turn it on anyway)
4     model.train()
5     # 1. Forward pass
6     y_pred = model(data)
7     # 2. Calculate the loss
8     loss = loss_fn(y_pred, labels)
9     # 3. Zero optimizer gradients
10    optimizer.zero_grad()
11    # 4. Perform backpropagation
12    loss.backward()
13    # 5. Perform gradient descent
14    optimizer.step()
15    return loss
```

```
1 # Test function
2 def test_step(model, loss_fn, data, labels):
3     # Turn on evaluation mode
4     model.eval()
5     # Setup inference mode context manager
6     with torch.inference_mode():
7         # 1. Forward pass
8         test_pred = model(data)
9         # 2. Calculate the loss
10        test_loss = loss_fn(test_pred, labels)
11        return test_loss
```

06 optimization.ipynb

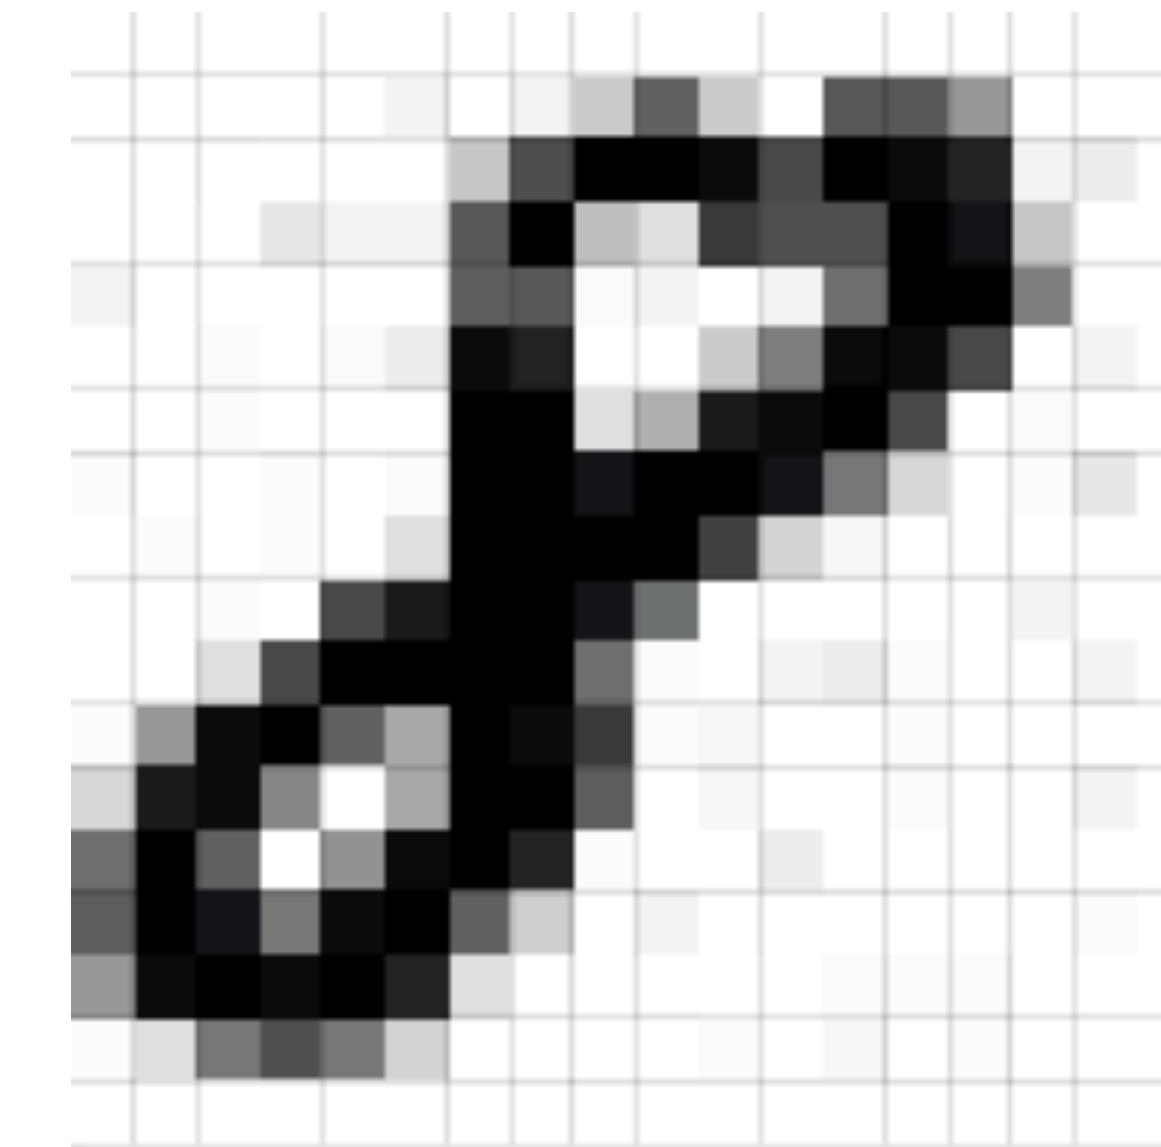
07 saving and loading models.ipynb

# Questions

# MNIST Example

<https://github.com/pytorch/examples/blob/main/mnist/main.py>

An Image is  
a matrix of  
pixel values



# Convolution

We slide the orange matrix over our original image (green) by 1 pixel (also called ‘stride’) and for every position, we compute element wise multiplication (between the two matrices) and add the multiplication outputs to get the final integer which forms a single element of the output matrix (pink). Note that the  $3 \times 3$  matrix “sees” only a part of the input image in each stride.

The primary purpose of Convolution in case of a ConvNet is to extract features from the input image.

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

1	0	1
0	1	0
1	0	1

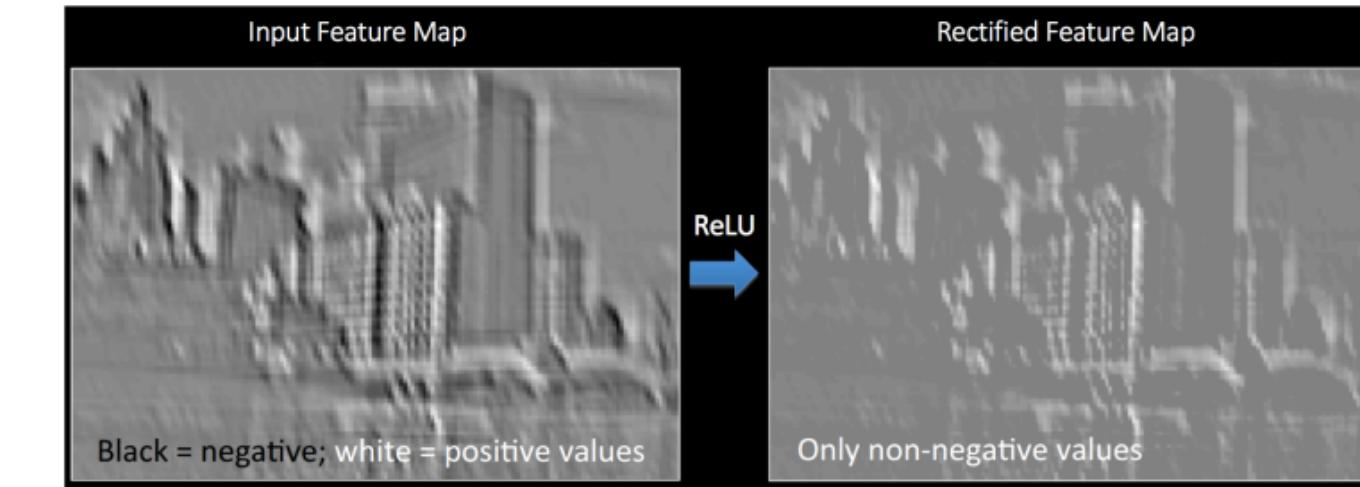
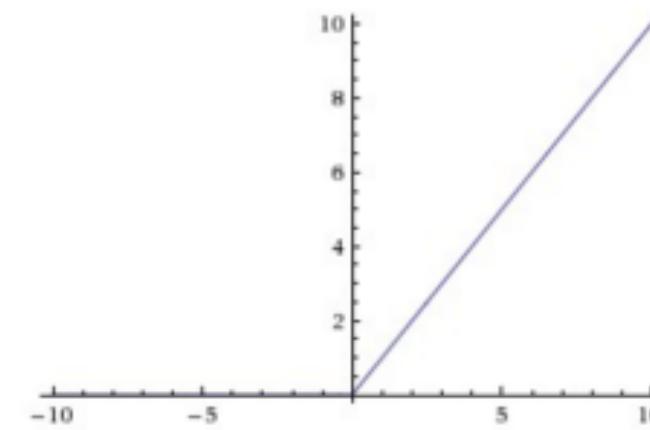
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image  
Convolved Feature

## Introducing Non Linearity (ReLU)

ReLU stands for Rectified Linear Unit and is a non-linear operation. Its output is given by:

Output = Max(zero, Input)



ReLU is an element wise operation (applied per pixel) and replaces all negative pixel values in the feature map by zero. The purpose of ReLU is to introduce non-linearity in our ConvNet, since most of the real-world data we would want our ConvNet to learn would be non-linear.

# Convolution

Different filters can give use different feature detectors.  
Note that convolution preserves spatial structure.



Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

# Pooling

Spatial Pooling (also called subsampling or downsampling) reduces the dimensionality of each feature map but retains the most important information. Spatial Pooling can be of different types: Max, Average, Sum etc.

**Example:** 2D Max Pooling with kernel size of 2. We slide our  $2 \times 2$  window by 2 cells (also called ‘stride’) and take the maximum value in each region.

